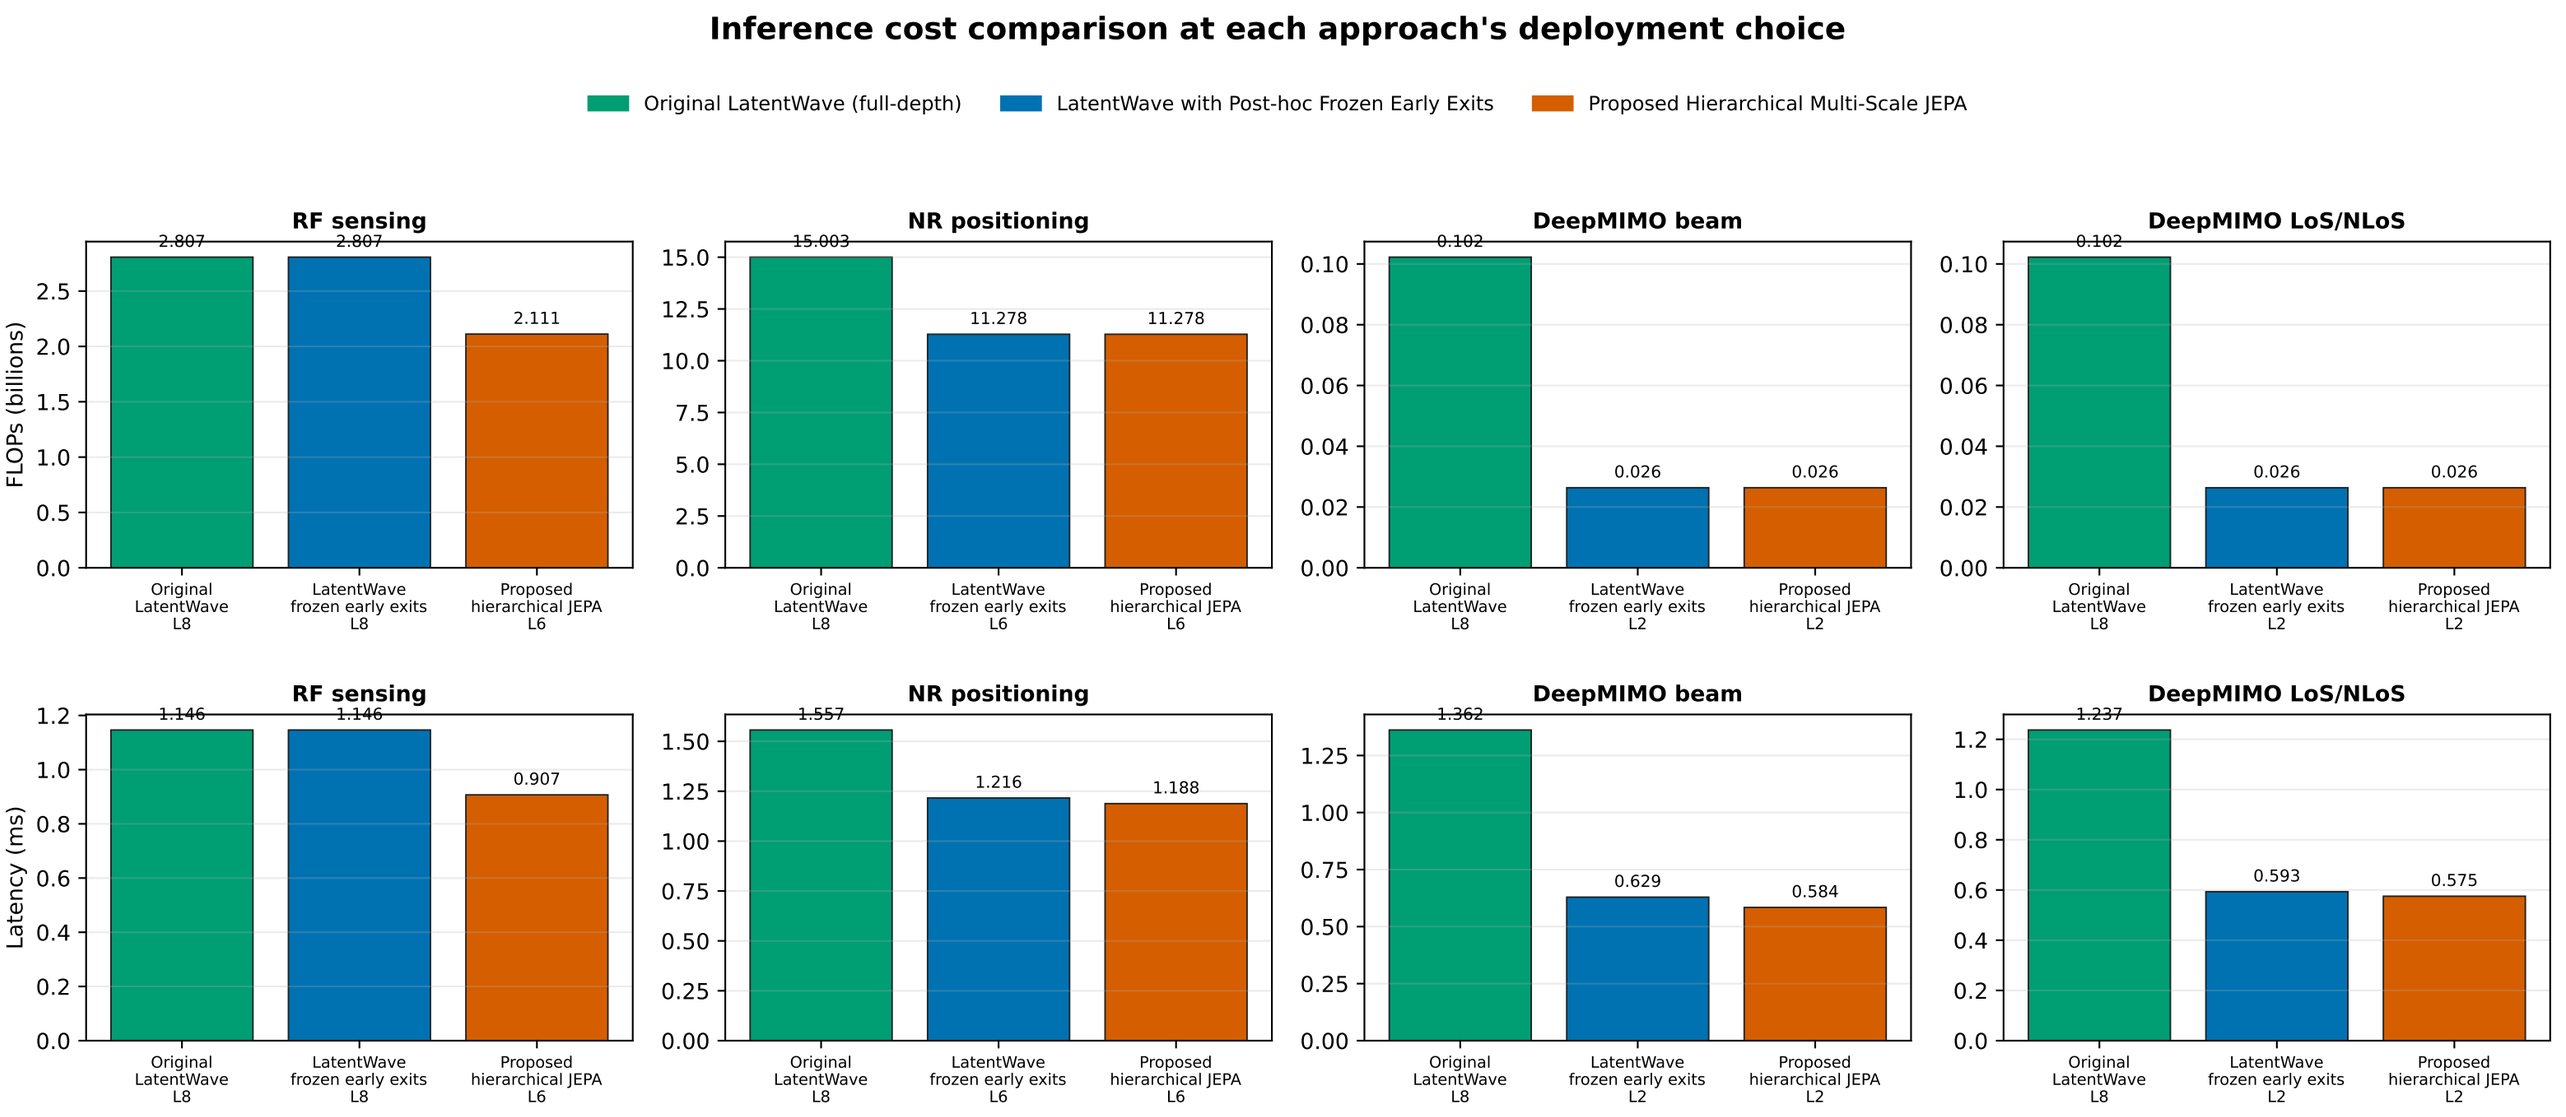
Inference cost comparison at each approach's deployment choice
Original LatentWave (full-depth)
LatentWave with Post-hoc Frozen Early Exits
Proposed Hierarchical Multi-Scale JEPA
RF sensing
NR positioning
DeepMIMO beam
DeepMIMO LoS/NLoS
FLOPs (billions)
2.807
2.807
2.111
15.003
11.278
11.278
0.102
0.026
0.026
0.102
0.026
0.026
Latency (ms)
1.146
1.146
0.907
1.557
1.216
1.188
1.362
0.629
0.584
1.237
0.593
0.575
Original LatentWave L8
LatentWave frozen early exits L8
LatentWave frozen early exits L6
LatentWave frozen early exits L2
Proposed hierarchical JEPA L6
Proposed hierarchical JEPA L2
Batch-1 synchronized CUDA timing on the RTX 5070 Ti; lower is better. Layer labels are printed below each method.